

Interpreting Credit-Risk Decisions Through a Diverse Counterfactual Explanation Framework

Anonymous Author(s)
Affiliation withheld for review

Abstract—Machine-learning credit-risk models support scalable assessment, but an adverse prediction alone offers little guidance about what would need to differ for the same model to produce a favorable outcome. This paper examines whether diverse counterfactual explanations remain available and interpretable when the predictive model is changed. Five classifier families—class-weighted logistic regression, decision tree, random forest, AdaBoost, and XGBoost—are trained on the German Credit, Give Me Some Credit, and Taiwan Credit datasets using identical stratified splits and preprocessing. For the ten test cases assigned the highest default probabilities by each classifier, the model-agnostic genetic-search implementation of Diverse Counterfactual Explanations (DiCE) requests four alternative profiles that reduce predicted default probability below 0.50. The search preserves designated immutable attributes and imposes dataset-specific directional restrictions. XGBoost achieved the highest ROC-AUC on all datasets (0.809, 0.869, and 0.777), while the decision tree achieved the lowest (0.649, 0.688, and 0.673). Counterfactual coverage was 100% in 13 of 15 model-dataset settings; logistic regression reached only 50% on Give Me Some Credit, and AdaBoost reached 90% on German Credit. The decision tree generally returned the sparsest explanations, but this simplicity coincided with weaker predictive discrimination. Taiwan Credit required substantially more simultaneous changes for every classifier. The results show that predictive performance and counterfactual recourse quality are distinct: conclusions obtained from one classifier may not generalize to another. These counterfactuals characterize model decision boundaries rather than causal interventions and require domain validation before use in lending.

Index Terms—credit scoring, counterfactual explanations, explainable artificial intelligence, algorithmic recourse, DiCE

I. INTRODUCTION

Machine-learning models are increasingly studied for credit scoring because they can extract predictive structure from heterogeneous applicant and repayment data. Their use in a high-impact domain nevertheless raises a practical interpretability problem. A probability of default communicates a model assessment, but it does not identify what would need to differ for the model to reach another outcome. Feature importance is insufficient for this purpose: an influential variable may be immutable, infeasible to change, or incapable of reversing a prediction by itself.

Counterfactual explanations address this gap by constructing alternative profiles that receive a desired model output. Rather than stating only that repayment history was influential, a counterfactual may specify a combination of lower utilization, fewer delinquencies, or different loan terms that moves a classifier below its adverse threshold. This form of explanation is local, contrastive, and linked directly to a decision [1], [2].

Useful explanations should, however, be valid, close to the original case, sparse, plausible, diverse, robust, and restricted to attributes that may legitimately vary [3]–[6].

The predictive model is itself an important source of variation. Linear, piecewise-constant, bagged, and boosted boundaries can assign different risks to the same population and require different changes to cross a common threshold. A study based on one classifier may therefore conflate properties of the explanation algorithm with properties of the underlying model.

This paper asks: *How does predictive model choice affect the availability and characteristics of genetically optimized counterfactual explanations in credit-risk classification?* We hold the explanation framework constant and compare five classifier families across three benchmark datasets. Our contributions are: (1) a controlled multi-model genetic-DiCE evaluation using common preprocessing and constraints; (2) joint reporting of predictive performance, validity, coverage, sparsity, proximity, and diversity; and (3) evidence that predictive discrimination and access to concise recourse do not necessarily align.

II. RELATED WORK

Wachter, Mittelstadt, and Russell described counterfactual explanations as nearby cases that obtain a different automated outcome without disclosing a model’s internal logic [1]. Algorithmic-recourse research similarly emphasizes a person’s ability to obtain a desired classification by altering actionable inputs. Ustun, Spangher, and Liu demonstrated that standard modeling choices can materially affect recourse in linear classification [2]. Model-agnostic optimization, graph-based search, multi-objective search, and generative methods subsequently broadened the available design space [7]–[11].

Mothilal, Sharma, and Tan proposed Diverse Counterfactual Explanations (DiCE), which generates sets of alternatives rather than one solution [3]. Its formulation balances outcome validity and proximity while encouraging diversity. Multiple alternatives matter because decision subjects may face different practical constraints. Reviews distinguish validity, proximity, sparsity, plausibility, diversity, actionability, causality, fairness, and robustness as separate dimensions of counterfactual quality [4]–[6]. Recent work also emphasizes that a model-changing profile is not necessarily a realizable intervention and that recourse itself can create risk or privacy exposure [12]–[14].

Credit scoring is a natural but difficult recourse setting. Applicant records combine immutable demographics, potentially actionable financial conditions, and historical variables that cannot literally be rewritten. Correlated monthly measures may also yield mathematically valid but operationally implausible profiles. Credit-scoring benchmarks additionally show that model ranking depends on discrimination, calibration, class imbalance, and data quality [15], [16]. This study therefore treats DiCE outputs as explanations of fitted classifiers, not causal prescriptions.

III. METHODOLOGY

A. Data and Preprocessing

Three public benchmarks are used. German Credit contains 1,000 applicants and credit, savings, employment, housing, and personal attributes. Give Me Some Credit contains 150,000 records with utilization, income, debt, credit-line, dependent, age, and delinquency variables. Taiwan Credit contains 30,000 records with credit limits, demographics, monthly bills, payments, repayment status, and subsequent default. The binary outcome equals one for default.

Each dataset is divided into stratified 75% training and 25% test partitions with random seed 42. Numeric features are median-imputed and standardized. Categorical features are mode-imputed and one-hot encoded. The held-out samples contain 250, 37,500, and 7,500 observations, with default rates of 0.300, 0.067, and 0.221, respectively.

B. Predictive Models

Five classifiers are trained separately: class-weighted logistic regression; a balanced decision tree with minimum leaf size five; a 300-tree balanced random forest [17] with minimum leaf size two; AdaBoost [18] with 200 estimators, learning rate 0.5, and balanced sample weights; and XGBoost [19] with 300 estimators, maximum depth four, learning rate 0.05, row and feature subsampling of 0.8, and class-ratio weighting. Predicted probabilities at or above 0.50 are assigned to default. Evaluation includes ROC–AUC, average precision, accuracy, balanced accuracy, F1, and Brier score.

C. Genetic DiCE

Let $f(x)$ be the estimated default probability for profile x . A counterfactual x' is valid when

$$f(x') < 0.50. \quad (1)$$

For each model–dataset pair, the ten test cases with the largest predicted default probabilities form a high-risk stress-test cohort. DiCE’s model-agnostic genetic search requests four class-0 alternatives per query. Proximity, diversity, and sparsity weights are 1.0, 1.0, and 0.5; post-hoc sparsity is 0.1. The objective balances prediction loss, distance, and sparsity while rewarding diversity among alternatives [3].

Dataset-specific policies fix immutable attributes. German Credit fixes age, personal-status/sex, and foreign-worker status; duration, credit amount, and installment rate may only

TABLE I
HELD-OUT ROC–AUC

Model	German	GMSC	Taiwan
Logistic regression	0.807	0.800	0.716
Decision tree	0.649	0.688	0.673
Random forest	0.797	0.851	0.768
AdaBoost	0.794	0.861	0.768
XGBoost	0.809	0.869	0.777

decrease. Give Me Some Credit fixes age; income may increase, while utilization, debt ratio, and delinquencies may decrease. Taiwan fixes sex, education, marital status, and age; credit limit and payment amounts may increase, while delay indicators may decrease.

Validity is independently recalculated from returned profiles. Coverage is the fraction of queries receiving at least one valid result. Sparsity is the number of original features changed. Proximity uses IQR-normalized numeric differences and unit categorical mismatch penalties:

$$d(x, x') = \frac{1}{p} \left[\sum_{j \in N} \frac{|x_j - x'_j|}{\text{IQR}(X_j)} + \sum_{j \in C} \mathbf{1}(x_j \neq x'_j) \right]. \quad (2)$$

Diversity is the mean pairwise mixed-type distance within a returned set. Raw distances are interpreted within datasets because feature spaces differ.

IV. RESULTS

A. Predictive Performance

Table I reports ROC–AUC. XGBoost ranked first on all datasets, while the single tree ranked last. Accuracy alone was misleading under imbalance: on Give Me Some Credit, random forest accuracy was 0.934 but balanced accuracy was 0.638, whereas XGBoost achieved 0.799 and 0.787, respectively.

B. Counterfactual Performance

Table II summarizes genetic DiCE. Thirteen of 15 settings achieved full coverage. Logistic regression covered only five of ten Give Me Some Credit cases; AdaBoost missed one German query. Most generated profiles were valid, although decision-tree and XGBoost validity on Give Me Some Credit were 92.5% and 95.0%.

The decision tree produced the sparsest explanations across all datasets, but also the weakest discrimination. Taiwan required many more changes under every classifier, reflecting coordinated alterations across monthly repayment, billing, and payment variables. Thus, representational complexity translated into larger counterfactual burden.

German counterfactuals were relatively close under all models. Give Me Some Credit yielded extreme distances for some models: mean logistic proximity was 53.32 and random-forest proximity 17.50. The data contain special delinquency values such as 96 and 98; treating them as ordinary counts can dominate normalized distance and create extreme scores. These

TABLE II
GENETIC DiCE RESULTS FOR TEN HIGH-RISK QUERIES PER CELL

Dataset	Model	Generated	Validity (%)	Coverage (%)	Mean changes
German	Logistic	35	100.0	100	4.29
	Tree	40	100.0	100	2.70
	Forest	32	100.0	100	3.03
	AdaBoost	31	100.0	90	3.29
	XGBoost	35	100.0	100	3.26
GMSD	Logistic	17	100.0	50	5.12
	Tree	40	92.5	100	4.00
	Forest	39	100.0	100	4.56
	AdaBoost	40	100.0	100	6.45
	XGBoost	40	95.0	100	5.29
Taiwan	Logistic	37	100.0	100	13.46
	Tree	39	100.0	100	10.18
	Forest	38	100.0	100	12.45
	AdaBoost	29	100.0	100	14.52
	XGBoost	38	100.0	100	13.84

magnitudes are therefore evidence of data-representation sensitivity, not intrinsically better diversity.

C. Applicant-Level Counterfactual Scenarios

Table III translates selected saved outputs into applicant-level scenarios. The XGBoost German case illustrates DiCE’s intended value: four distinct profiles moved the same prediction from 0.972 to between 0.141 and 0.247. Alternatives traded shorter duration, lower credit amount, savings category, housing, and guarantor status. The explanation set therefore exposes different routes through the model boundary instead of presenting one recommendation as uniquely necessary.

The latter two scenarios demonstrate why threshold validity is not equivalent to useful recourse. The Give Me Some Credit alternative reduces special delinquency codes but leaves values that are not credible ordinary counts. The Taiwan alternative modifies historical monthly records rather than a prospective action. We retain these cases because they reveal model and data defects that a counterfactual audit should surface, but we do not present them as recommendations.

V. DISCUSSION

The results demonstrate that counterfactual explanations are conditional on the model being explained. Holding the generator, threshold, and constraints constant did not eliminate differences in coverage or burden. A single-model study may consequently mistake decision geometry for a general property of credit recourse. This observation complements earlier findings that model development affects recourse availability [2], [20].

No classifier dominated both prediction and explanation. XGBoost offered the strongest ROC–AUC, whereas the decision tree offered the fewest changes. The latter simplicity should not be interpreted as superiority because it coincided with materially weaker prediction. Coverage also matters: explanations that are excellent for easy cases but unavailable to extreme-risk cases cannot provide universal recourse.

Validity alone remains insufficient. A classifier changes output when supplied a returned profile, but the profile may be distant, historically impossible, or unavailable to an applicant. These explanations are algorithmic rather than causal: they do not show that enacting a change would reduce real default risk.

VI. LIMITATIONS

The study is a ten-case-per-model stress test, not a population estimate. Each model selects its own highest-risk cases, preventing paired applicant-level inference. Repeated seeds are required to quantify genetic-search stability. Constraints do not fully encode temporal or institutional feasibility. Special delinquency codes require cleaning and sensitivity analysis. Distances are not directly comparable across heterogeneous datasets. Freezing protected attributes is not a fairness analysis, and subgroup disparities in coverage and cost remain untested. Finally, the datasets contain default outcomes rather than historical lender decisions; “adverse” denotes predicted default, not observed denial.

VII. CONCLUSION AND FUTURE WORK

Across five classifiers and three credit datasets, model choice materially altered genetic-DiCE coverage and explanation complexity. XGBoost achieved the best predictive discrimination, while a single tree produced the sparsest explanations at a predictive cost. Logistic regression failed to provide recourse for half of the most extreme Give Me Some Credit cases, and Taiwan required many simultaneous changes regardless of model.

Future work will evaluate a larger shared query cohort, repeat search across seeds, report confidence intervals and paired tests, clean anomalous codes, distinguish strict applicant recourse from permissive model auditing, and assess subgroup disparities. These steps are necessary before counterfactual explanations can be treated as equitable guidance rather than descriptions of a model boundary.




Fig. 1. Predictive performance and genetic-DiCE recourse across model families.




Fig. 2. XGBoost probability shifts for high-risk German Credit queries. Each line joins an original prediction to a generated profile; the dashed line is the 0.50 threshold.




Fig. 3. Features most frequently changed by valid XGBoost counterfactuals in Taiwan Credit. Repeated monthly measures explain the comparatively low sparsity.

REFERENCES

- [1] S. Wachter, B. Mittelstadt, and C. Russell, “Counterfactual explanations without opening the black box: Automated decisions and the gdpr,” *Harvard Journal of Law & Technology*, vol. 31, no. 2, pp. 841–887, 2018.
- [2] B. Ustun, A. Spangher, and Y. Liu, “Actionable recourse in linear classification,” in *Proceedings of the Conference on Fairness, Accountability, and Transparency*, 2019, pp. 10–19.
- [3] R. K. Mothilal, A. Sharma, and C. Tan, “Explaining machine learning classifiers through diverse counterfactual explanations,” in *Proceedings*

TABLE III
ILLUSTRATIVE COUNTERFACTUAL SCENARIOS RECONSTRUCTED FROM SAVED GENETIC-DICE OUTPUTS

Dataset/model	Default probability	CF probability	Changes from original profile
German/XGB CF1	0.972	0.234	Duration 48→14 months; credit amount 4308→338; savings A61→A64.
German/XGB CF2	0.972	0.247	Duration 48→24; credit amount 4308→3777; savings A61→A64; housing A151→A152.
German/XGB CF3	0.972	0.195	Duration 48→14; credit amount 4308→3973; savings A61→A64; housing A151→A153.
German/XGB CF4	0.972	0.141	Duration 48→4; credit amount 4308→338; debtor/guarantor A101→A103; housing A151→A153.
GMSR/RF	0.989	0.418	Utilization 1.00→0; debt ratio 22→0; delinquency codes 98→89, 81, and 53 across the three delinquency variables.
Taiwan/tree	1.000	0.000	Six historical bill amounts and PAY_AMT4 changed; seven simultaneous changes in total.

of the 2020 Conference on Fairness, Accountability, and Transparency, 2020, pp. 607–617.

- [4] S. Verma, V. Boonsanong, M. Hoang, K. E. Hines, J. P. Dickerson, and C. Shah, “Counterfactual explanations and algorithmic recourses for machine learning: A review,” *arXiv preprint arXiv:2010.10596*, 2020.
- [5] M. T. Keane, E. M. Kenny, E. Delaney, and B. Smyth, “If only we had better counterfactual explanations: Five key deficits to rectify in the evaluation of counterfactual xai techniques,” in *Proceedings of the Thirtieth International Joint Conference on Artificial Intelligence*, 2021, pp. 4466–4474.
- [6] J. Jiang, F. Leofante, A. Rago, and F. Toni, “Robust counterfactual explanations in machine learning: A survey,” *arXiv preprint arXiv:2402.01928*, 2024.
- [7] A.-H. Karimi, G. Barthe, B. Balle, and I. Valera, “Model-agnostic counterfactual explanations for consequential decisions,” in *Proceedings of AISTATS*, ser. PMLR, vol. 108, 2020, pp. 895–905.
- [8] R. Poyiadzi, K. Sokol, R. Santos-Rodriguez, T. D. Bie, and P. Flach, “Face: Feasible and actionable counterfactual explanations,” in *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, 2020, pp. 344–350.
- [9] C. Russell, “Efficient search for diverse coherent explanations,” in *Proceedings of the Conference on Fairness, Accountability, and Transparency*, 2019, pp. 20–28.
- [10] S. Dandl, C. Molnar, M. Binder, and B. Bischl, “Multi-objective counterfactual explanations,” in *Parallel Problem Solving from Nature*, 2020, pp. 448–469.
- [11] D. Nemirovsky, N. Thiebaud, Y. Xu, and A. Gupta, “Counterfactual: Generating counterfactuals for real-time recourse and interpretability using residual gans,” in *Proceedings of UAI*, ser. PMLR, vol. 180, 2022, pp. 1488–1497.
- [12] A.-H. Karimi, B. Schölkopf, and I. Valera, “Algorithmic recourse: From counterfactual explanations to interventions,” in *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency*, 2021, pp. 353–362.
- [13] H. Fokkema, D. Garreau, and T. van Erven, “The risks of recourse in binary classification,” in *Proceedings of AISTATS*, ser. PMLR, vol. 238, 2024, pp. 550–558.
- [14] M. Pawelczyk, H. Lakkaraju, and S. Neel, “On the privacy risks of algorithmic recourse,” in *Proceedings of AISTATS*, ser. PMLR, vol. 206, 2023.
- [15] S. Lessmann, B. Baesens, H.-V. Seow, and L. C. Thomas, “Benchmarking state-of-the-art classification algorithms for credit scoring: An update of research,” *European Journal of Operational Research*, vol. 247, no. 1, pp. 124–136, 2015.
- [16] I.-C. Yeh and C. hui Lien, “The comparisons of data mining techniques for the predictive accuracy of probability of default of credit card clients,” *Expert Systems with Applications*, vol. 36, no. 2, pp. 2473–2480, 2009.
- [17] L. Breiman, “Random forests,” *Machine Learning*, vol. 45, pp. 5–32, 2001.
- [18] Y. Freund and R. E. Schapire, “A decision-theoretic generalization of on-line learning and an application to boosting,” *Journal of Computer and System Sciences*, vol. 55, no. 1, pp. 119–139, 1997.
- [19] T. Chen and C. Guestrin, “Xgboost: A scalable tree boosting system,” in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 785–794.
- [20] K. Rawal and H. Lakkaraju, “Beyond individualized recourse: Interpretable and interactive summaries of actionable recourses,” in *Advances in Neural Information Processing Systems*, vol. 33, 2020.